



Conception Avancée de Bases de Données



Google
Big Table



Traduction en cours



1996 1997



- Google started as a research project at Stanford University
- Created by Ph.D. candidates Larry Page and Sergey Brin when they were 24 years old and 23 years old respectively.
- BackRub became Google from « googol » 1 leading 100 zeros.

Time line



- 1998 : Susan Wojcicki garage
 - 232 Santa Margarita à Menlo Park.
- 2001 :
 - Eric Schmidt CEO
- 2004 : "Googleplex »
 - 1600 Amphitheatre Parkway à Mountain View
- August 2004 :
 - Wall street NASDAQ 19 605 052 shares 85 \$ (1000 \$).



Alphabet Inc Class A

NASDAQ: GOOGL - 2 déc. à 08:53 UTC-5

764,33 USD **0,00** (0,00 %)

Avant l'ouverture: 761,00 **+0,44 %**

1 jour

5 jours


1 mois

3 mois

1 an

5 ans

max



Ouverture

-

+Haut

-

+Bas

-

Capitalis.

513,38 Md

Cours/bén.

27,47

Rend. div.

-

Alphabet Inc Class C

NASDAQ: GOOG

[+ Follow](#)

[Overview](#)

[News](#)

[Compare](#)

[Financials](#)



1,048.65 USD +8.04 (0.77%) ↑

Dec 14, 3:31 PM EST · Disclaimer

[1 day](#)

[5 days](#)

[1 month](#)

[1 year](#)

[5 years](#)

[Max](#)



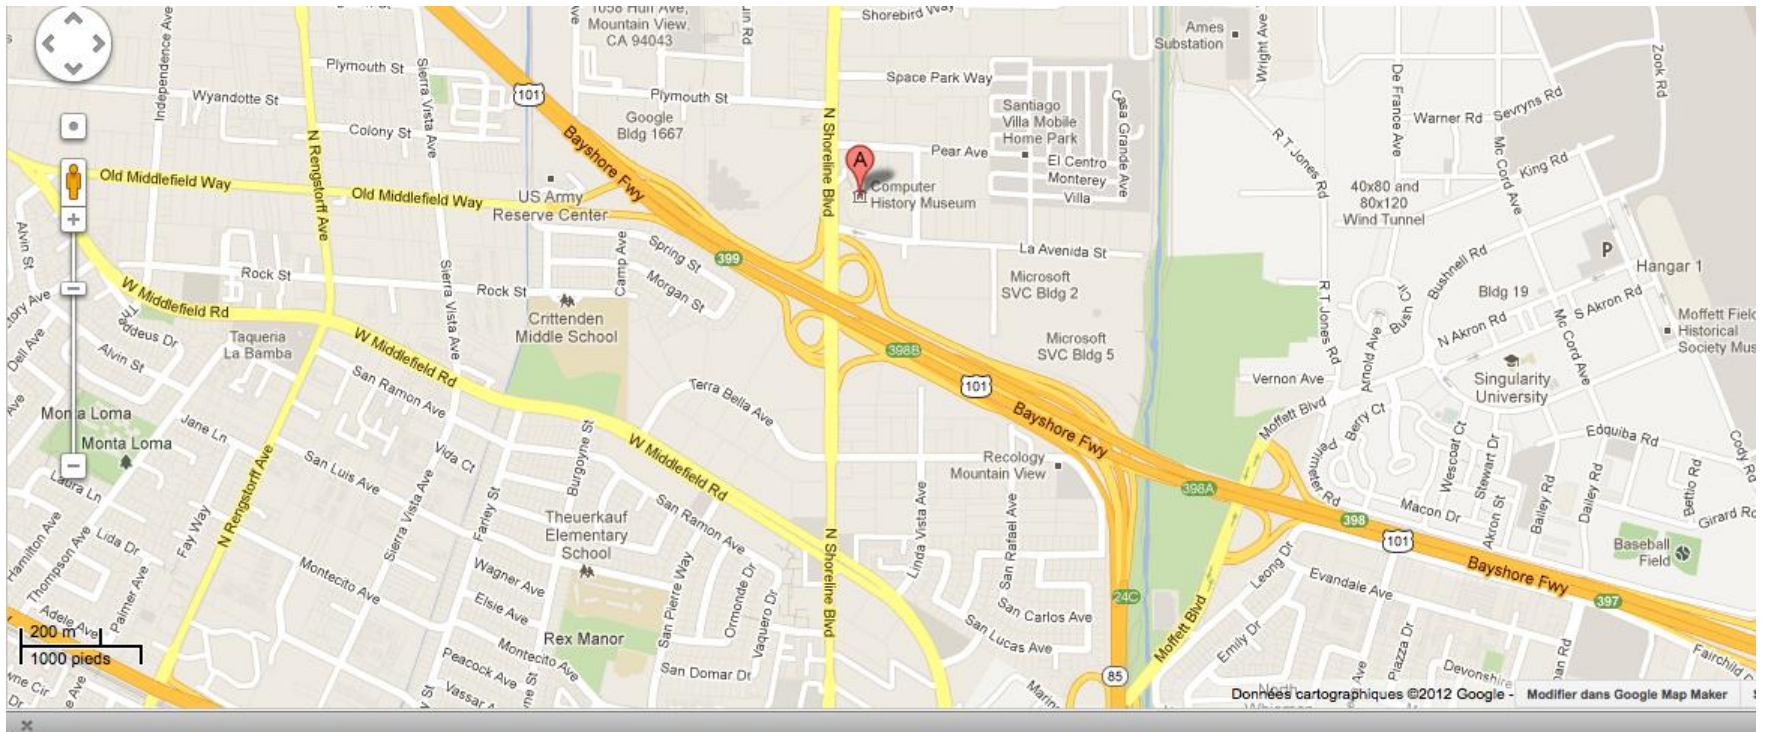
Open	1,045.00
High	1,058.50
Low	1,043.11
Mkt cap	732.41B
P/E ratio	34.94

Div yield	-
Prev close	1,040.61
52-wk high	1,062.38
52-wk low	770.41

Computer History Museum



- 1401 North Shoreline Boulevard Mountain View, CA 94043, United States



Sergey Brin & Larry Page



Susan Wojcicki : Google VP





Google Big Table



<http://actu.abondance.com/2008/04/google-pourrait-proposer-big-table-son.html>

Bigtable



- Bigtable is a distributed storage system for managing structured data that is designed to scale to a very large size: petabytes of data across thousands of commodity servers.
- Many projects at Google store data in Bigtable, including web indexing, Google Earth, and Google Finance.



Google Servers architecture





- PC mother boards and disks

Google rack



Low cost PC



FRYS.COM



FRYS Mountain View, CA 94043



Fry's Electronics, the first electronic superstore: 1177 Kern Ave, Sunnyvale









http://www.yelp.com/biz_photos/CsFTB9cErCkOkOqWGIqzHA?select=4RHMVug11I7cMSAaHHRLuA#4RHMVug11I7cMSAaHHRLuA





Google Story





Google Story




Google Story



Google Rack



Horizontal Scalability : Shared Disk



low-cost low performance servers



- Comodity computer
- Many low-cost low performance servers.
- Servers are aggregated in clusters with more than a thousand servers each
- A cluster consists of tens of racks.
- The nodes inside a rack (about 80 per rack) are connected by a switch with something like a 1gb/s data rate between any two nodes.
- High degree of parallelism

Bigtable



- A Bigtable is a sparse, distributed, persistent multidimensional sorted map.
- The map is indexed by a row key, column key, and a timestamp; each value in the map is an uninterpreted array of bytes.

KEY

VALUE

(row:string, column:string, time:int64) → string

Google Search




- The key is every possible search term.
- The value is the table of web sites to return for this search term.
- Index is replaced by an algorithm called MapReduce
- MapReduce is a parallel algorithm, the search can be done by multiple servers in parallel.

Hash Table



- Sparse 3D hash table :
 - row names,
 - column names,
 - versions (timestamps).
- Everything's a string (sequences of characters) :
 - row names,
 - column names,
 - data items

Keys value



Sort of 1 PB (10^{15}) in 6 hours



- One petabyte is a thousand terabytes,
- 10 trillion (10^{12}) 100-byte records
- on 4,000 computers
- It took six hours and two minutes



Refroidissement Google



