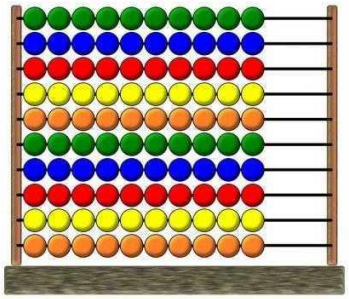

Lambda architectures

Probabilistic Data Structures

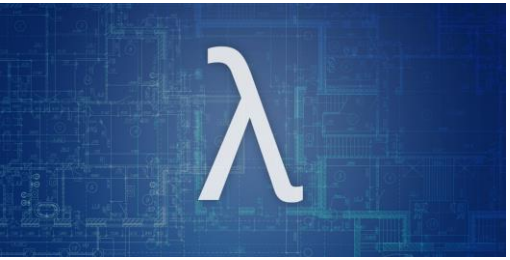

λ vs $\mathcal{U}$

Big Data Definition

- Big Data is a collection of large datasets that cannot be processed using traditional computing techniques.

- Any data that can challenge our current technology

Big Data

Big data is about the Insight that we want to extract from information.

Data

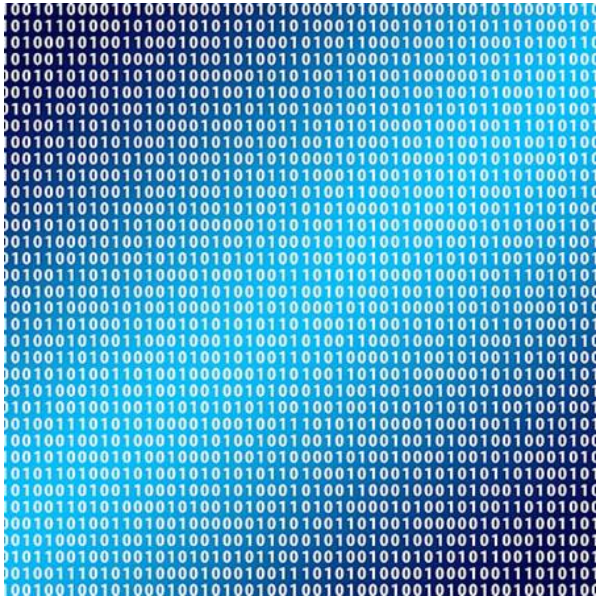

Computing

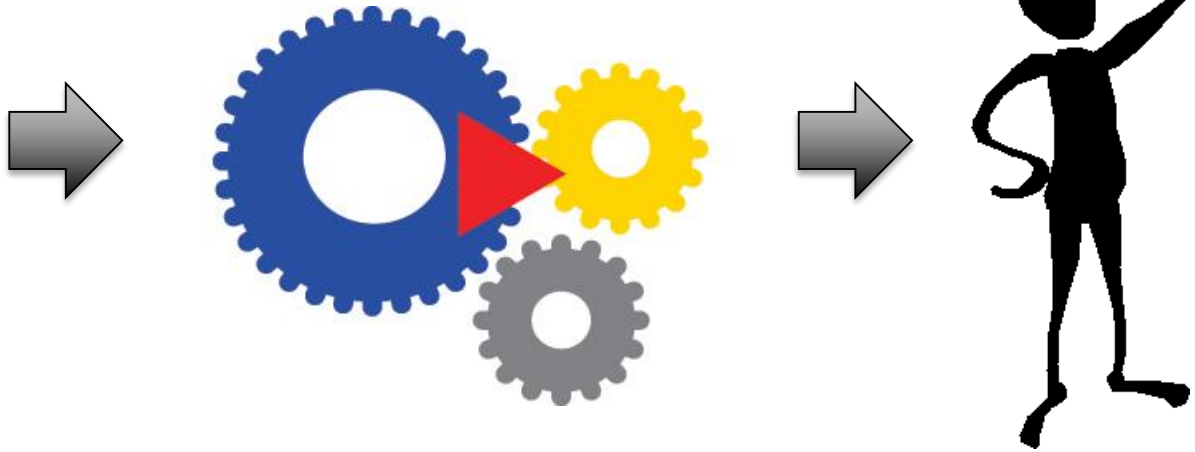

Lambda Architecture

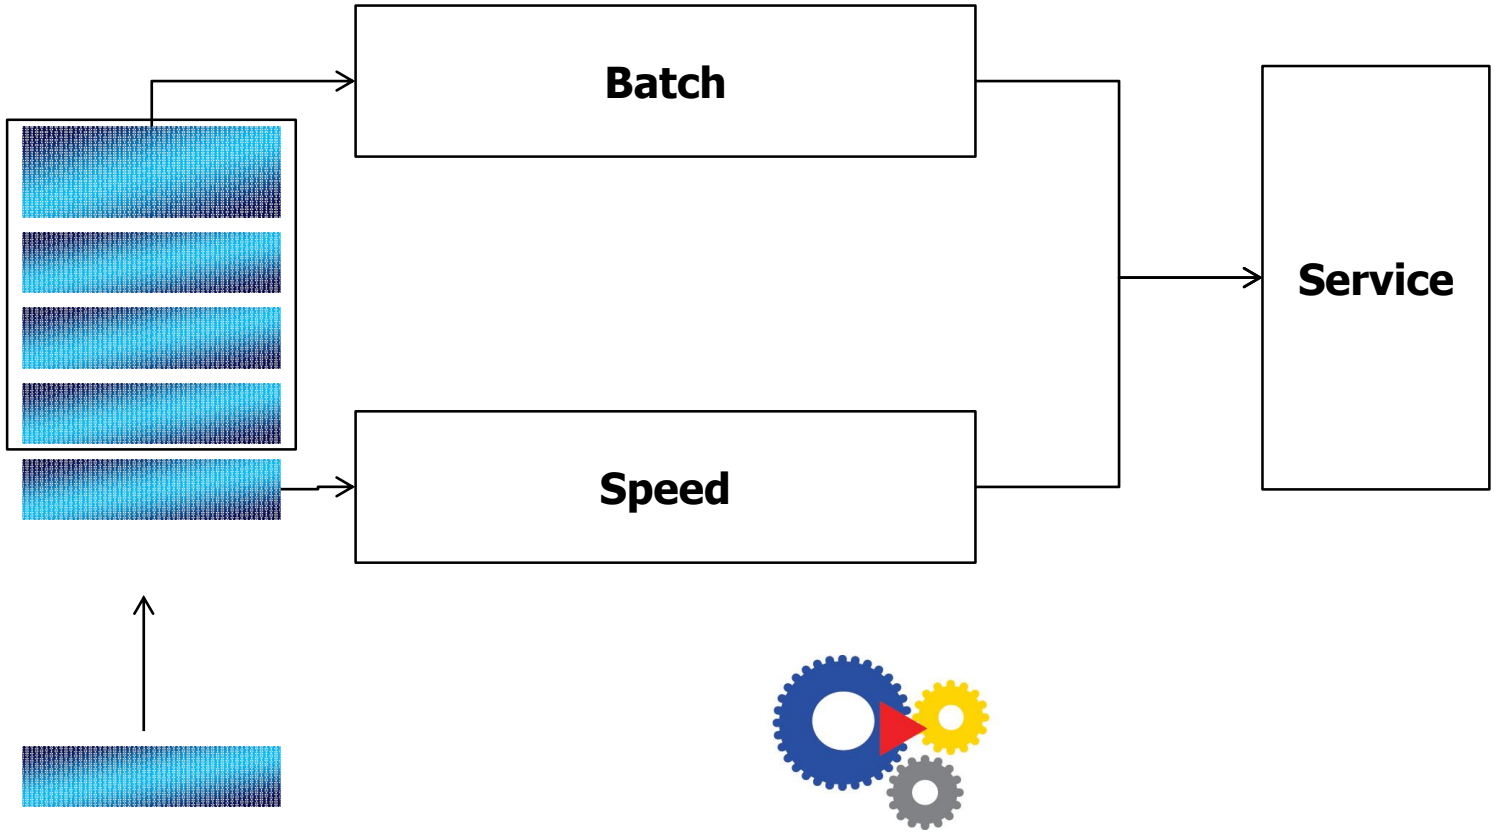

Computing

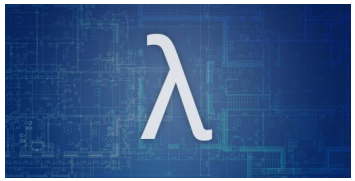

Lambda Architecture

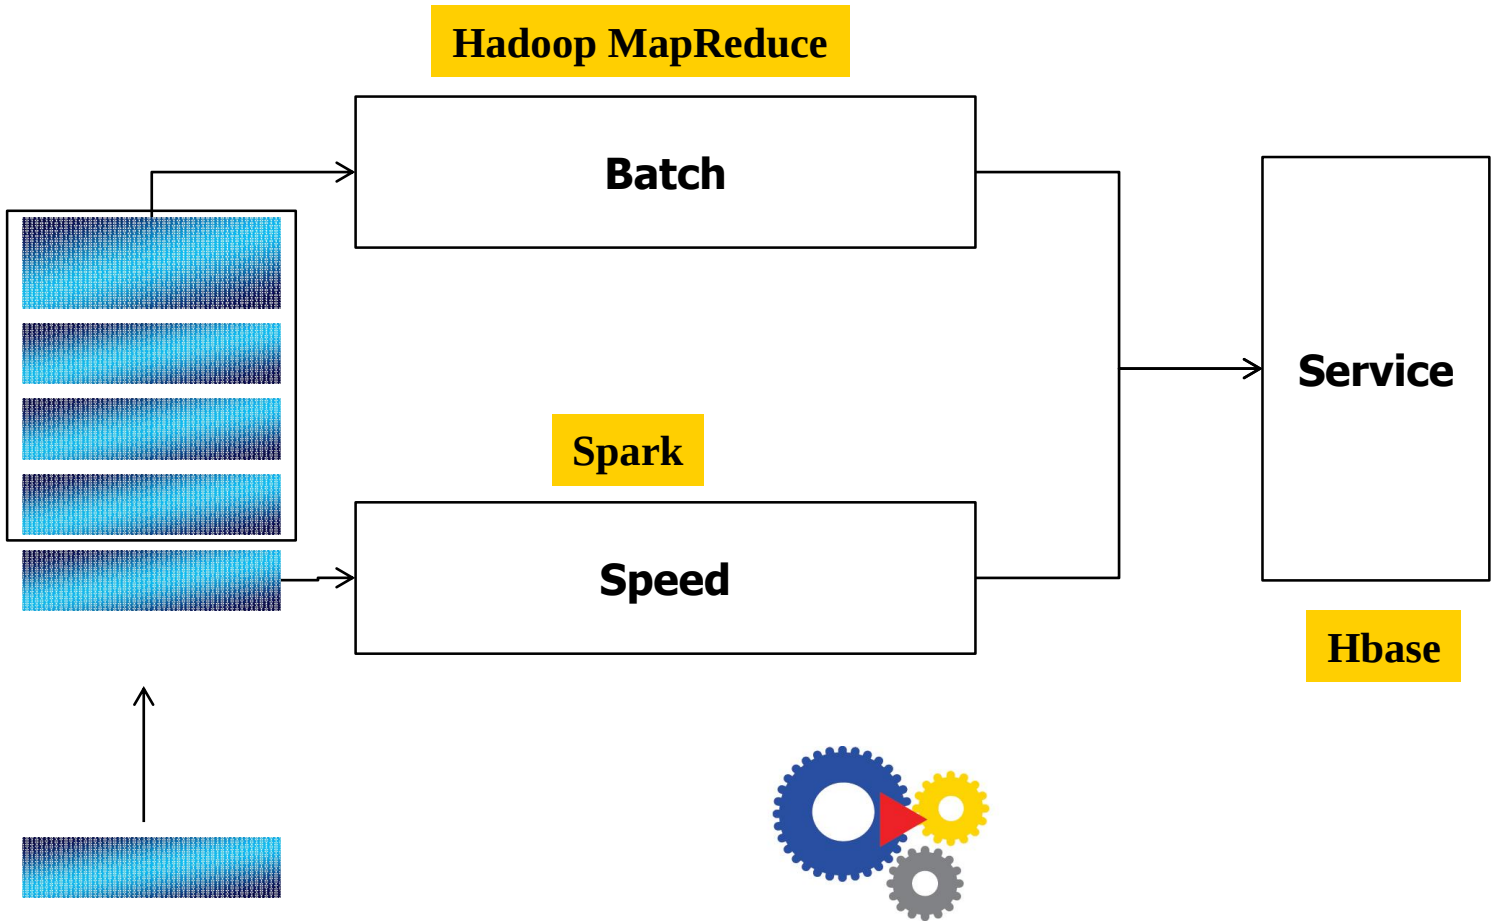

Lambda Architecture

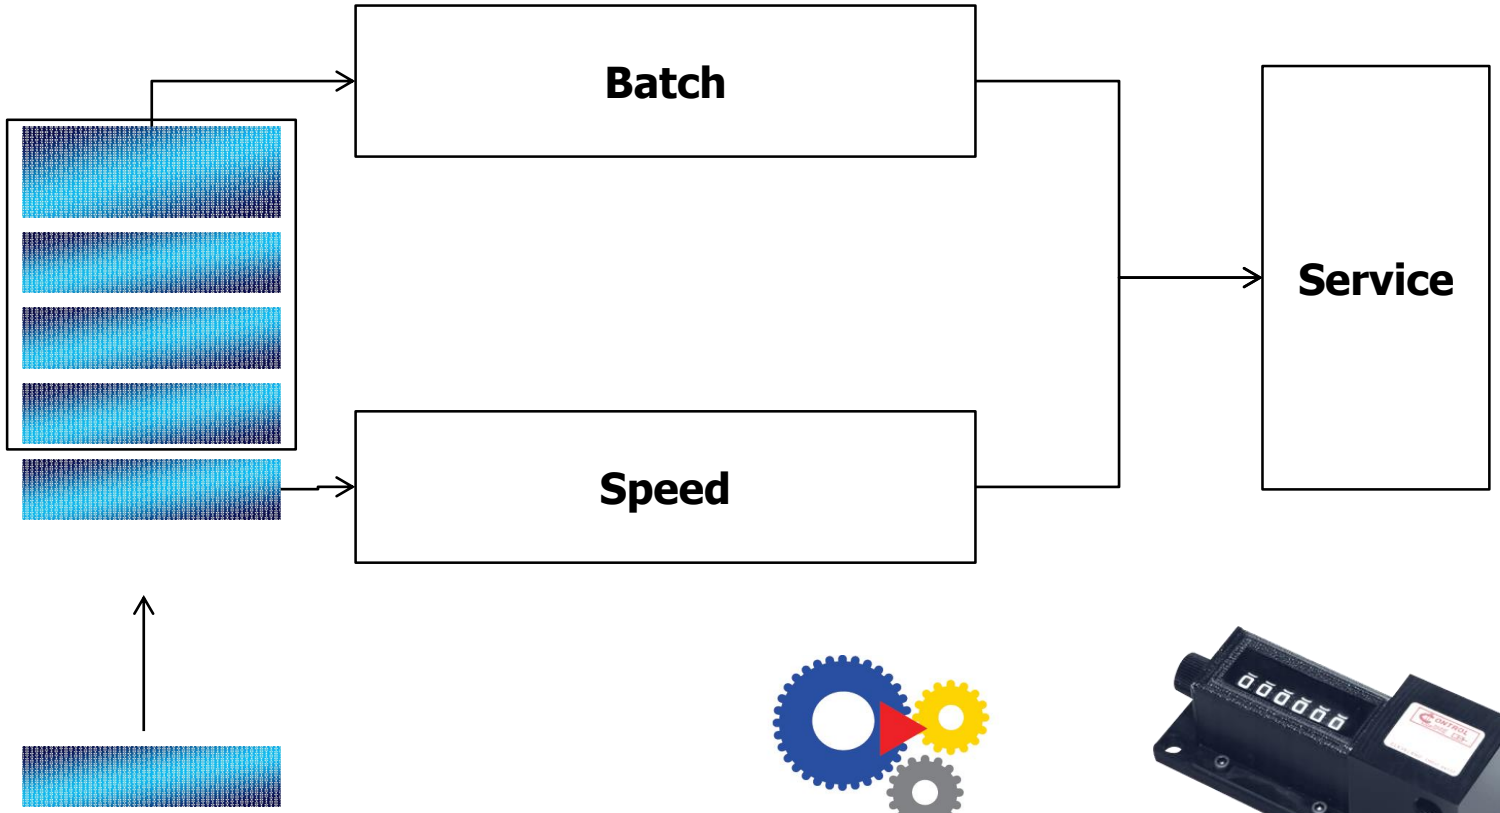

Computing

Lambda architecture

- 1) Speed/event layer :
  - Stream processing architecture.
  - The data are appended to the master data periodically over a time window.
  - Real time views
- 2) Batch layer :
  - Manages the master dataset:
    - Immutable, append-only raw data and pre-computing batch views.
  - Batch views
- 3) Serving layer:
  - Merges batch and real-time views.

Lambda architecture

- Offline or batch processing : insight
  - No time constraint : no data structures
  - Data are stored on disk
  - Data Lakes
- Streaming data : near real-time processing
  - Collecting analytics
  - Monitoring
  - Alerting,
  - Updating search indexes
  - Caches.

Data Stream

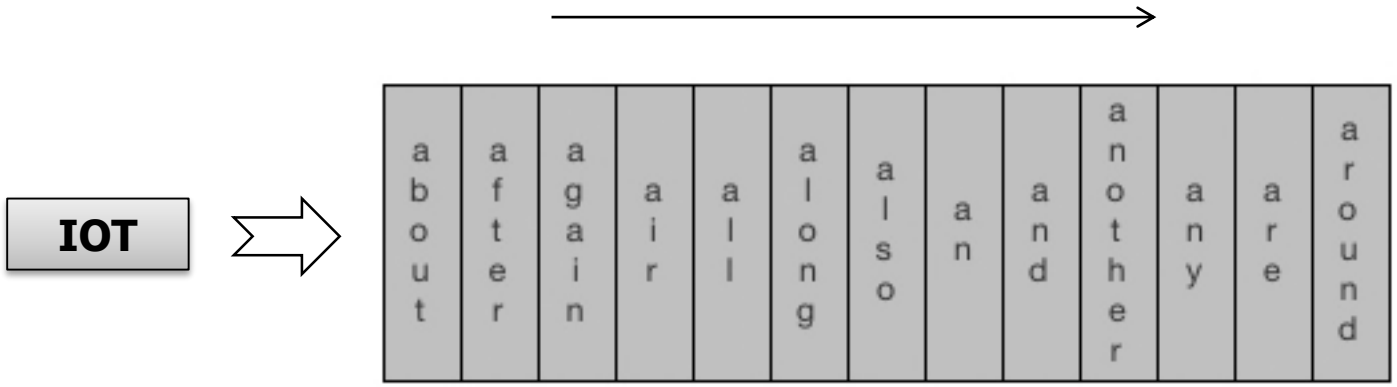
IOT
about
after
again
air
all
along
also
an
and
another
any
are
around

Data Stream

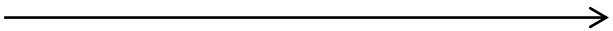

IOT

about
after
again
air
all
along
also
an
and
another
any
are
around

Data Stream

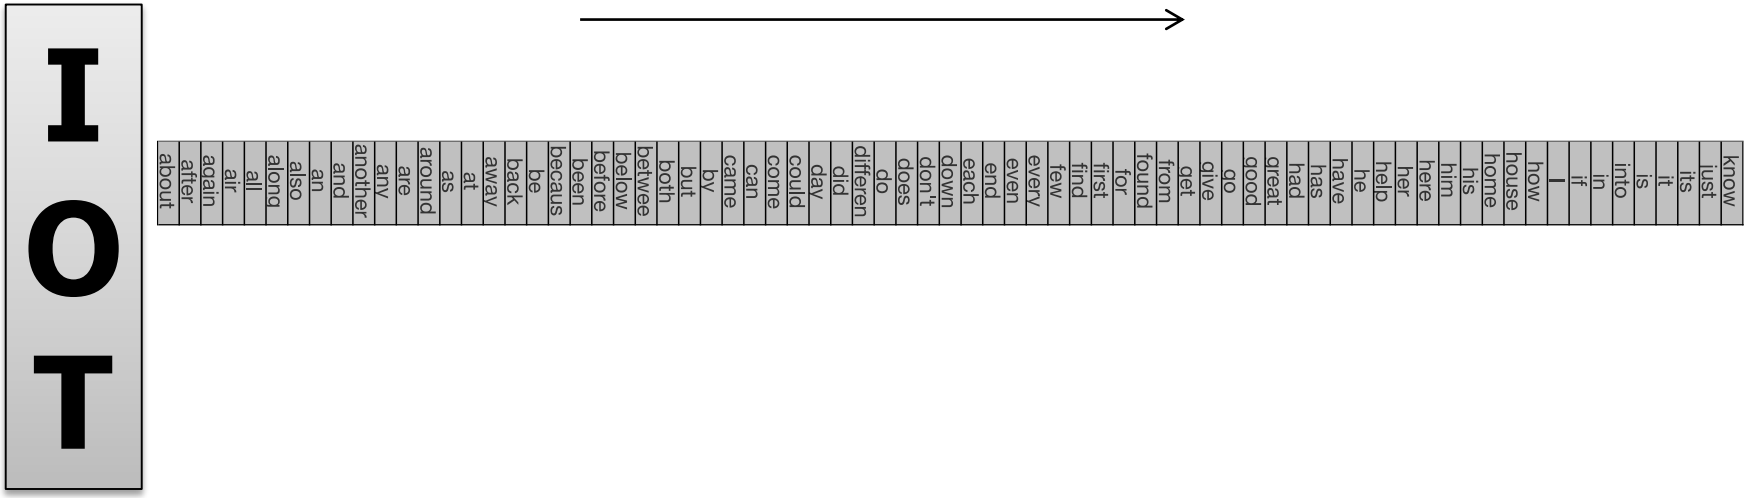
I
O
T
about
after
again
air
all
along
also
an
and
another
any
are
around
as
at
away
back
be
becaus
been
before
below
betwee
both
but
by
came
can
come
could
day
did
differen
do
does
don't
down
each
end
even
every
few
find
first
for
found
from
get
give
go
good
great
had
has
have
he
help
her
here
him
his
home
house
how
I
if
in
into
is
it
its
just
know

Data Stream

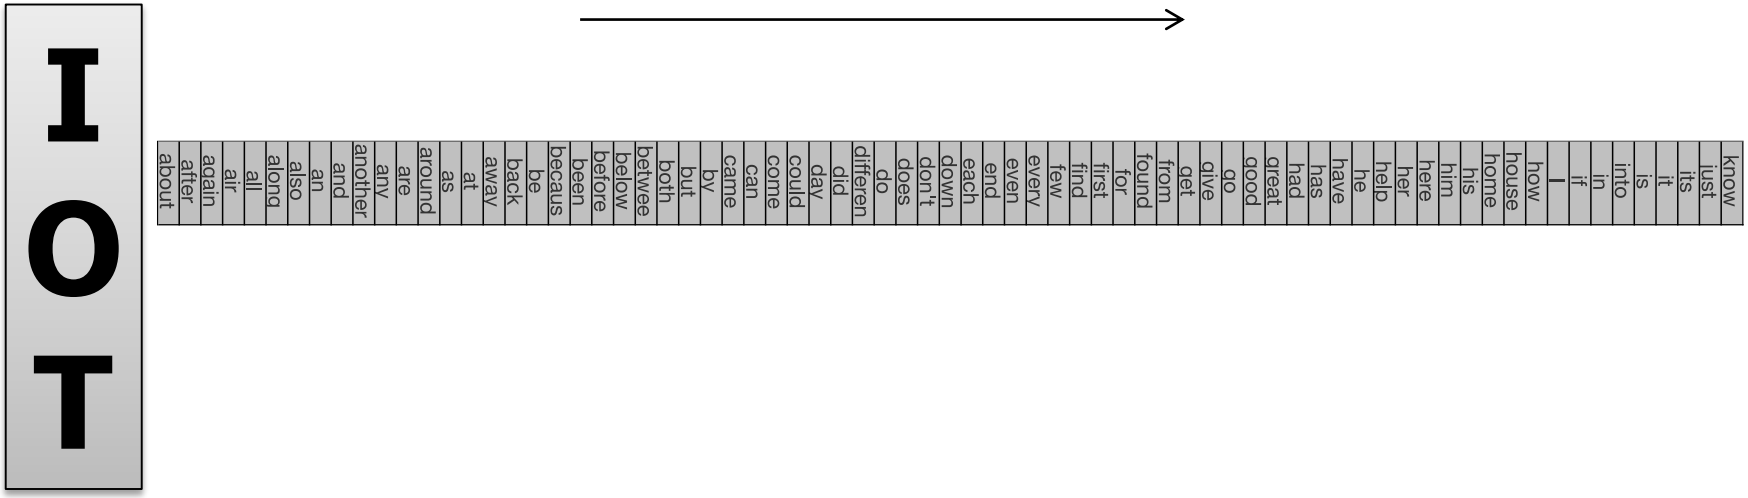
I
O
T
about
after
again
air
all
along
also
an
and
another
any
are
around
as
at
away
back
be
becaus
been
before
below
betwee
both
but
by
came
can
come
could
day
did
differen
do
does
don't
down
each
end
even
every
few
find
first
for
found
from
get
give
go
good
great
had
has
have
he
help
her
here
him
his
home
house
how
I
if
in
into
is
it
its
just
know

Data Stream

I O T

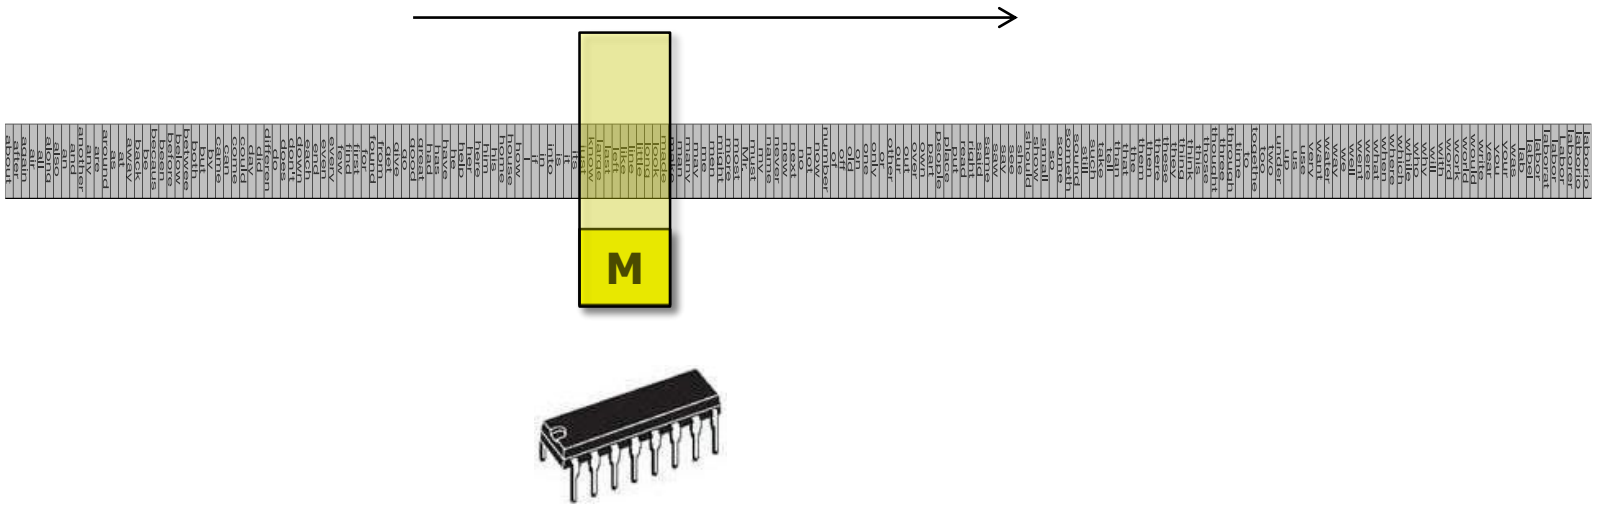

Data Stream

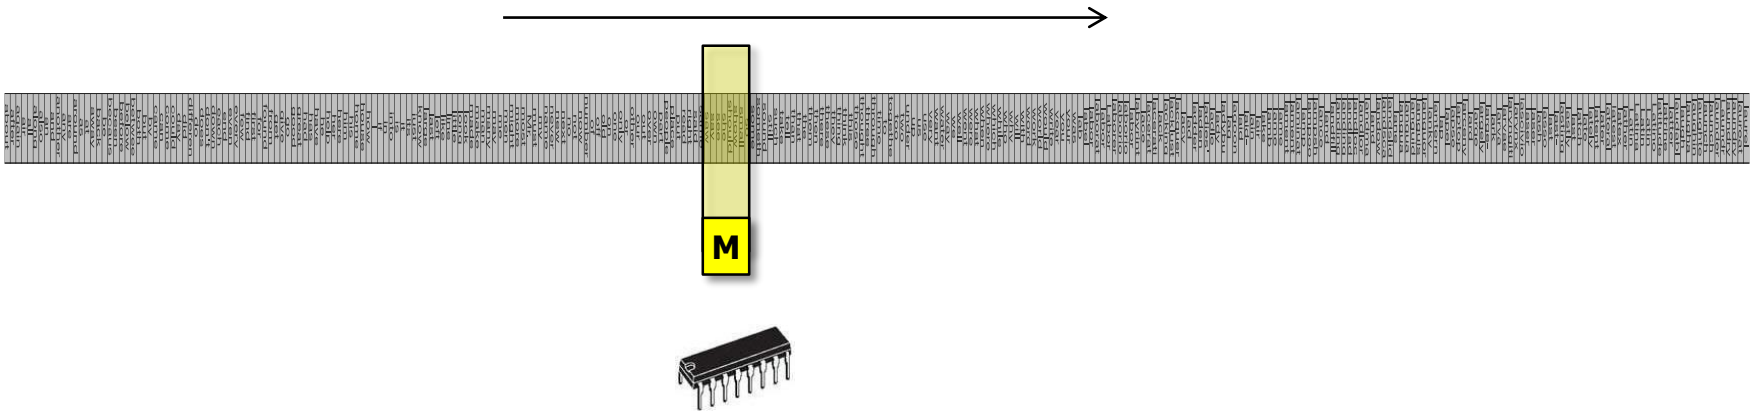

Data Stream

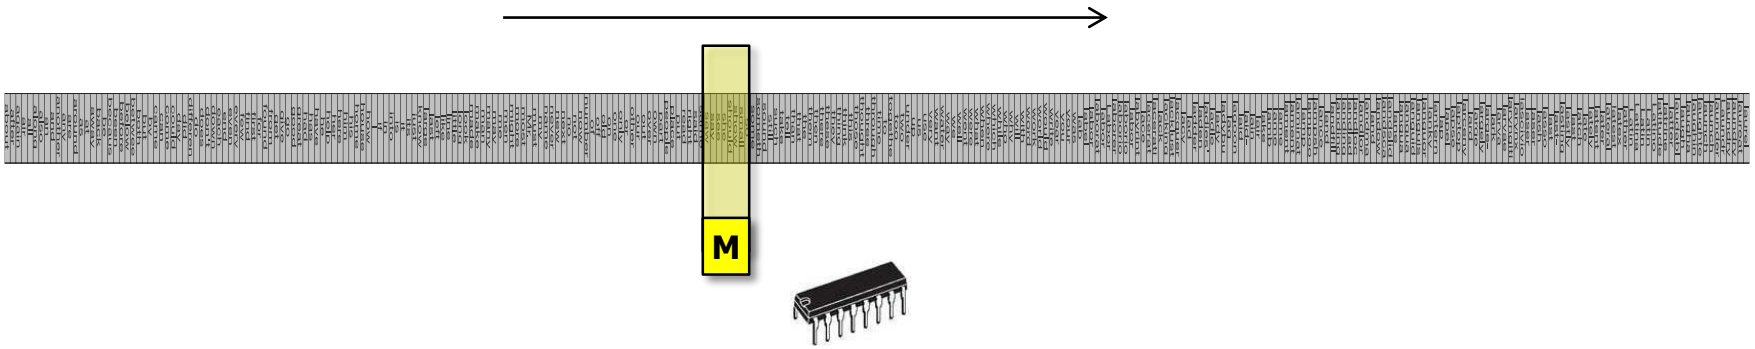

Who ?
How Many ?
What ?
Together ?
How many of them ?
Why ?

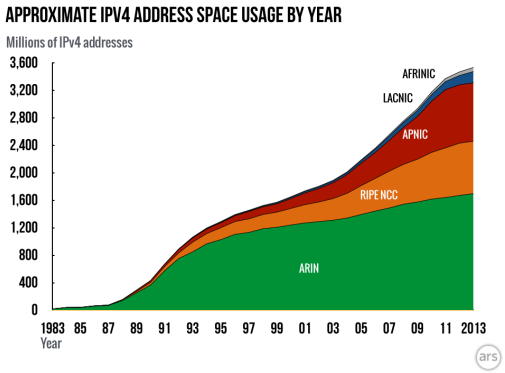

Speed/event layer

- Online, real-time processing
- Collecting real-time analytics
- In stream processing
- Large (Big) data streams :
  - Impractical to store that data in memory
  - Trade off to get real-time insight

- Stream Processing with Probabilistic counting
  - In-stream counters

Advantages Probabilistic Data Structures

- Use small amount of memory
  - Can be controled
- Can be easily parallelizable
  - Hashes are independent
- Have constant query time
  - Not even amortized constant like in dictionary

Usage

- Computation Of Advanced Metrics :
  - Number Of Unique Visitor
  - Most Frequent Items
- Web Crawlers:
  - Process A Stream Of Urls And Documents Tp Produce Indexed Content.
- Websites:
  - Process A Stream Of Page Views And Update Various Counters And Gauges.
- Email:
  - Process A Stream Of Text And Produce A Filtered, Spam-free Inbox.

Exemple

- How to count the cardinality of IP addresses which access our website ?
- There are over four billion possible distinct IP V4 addresses.

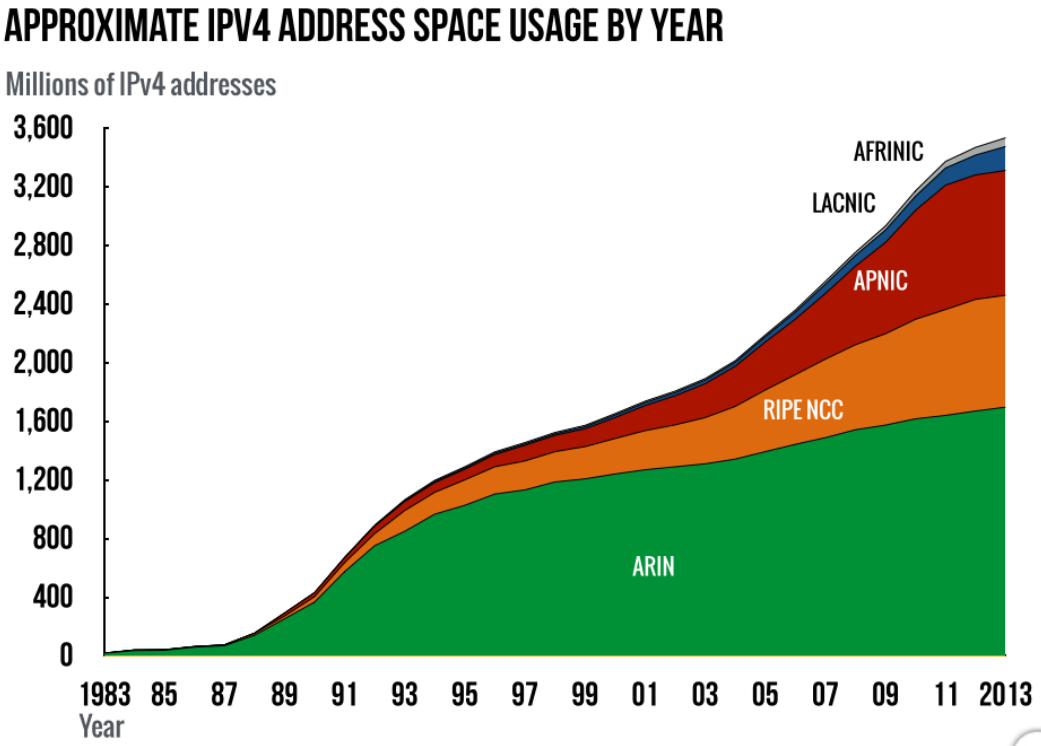

Probabilistic data structures

- Probabilistic data structures can not provide definite answer
- Instead they provide a reasonable approximation of the answer and a way to approximate this estimation.
- Useful for big data and streaming application because they allow to decrease the amount of memory needed with comparison to data structures that give you exact answers.

Membership

Is this guy on board ?

How many of them are there ?

Presence evaluation

- How to evaluate the presence of the element in a stream ?

Frequency evaluation

- How to store a numerical value associated with each element ?
- How to evaluate the number of occurrences of the element in a stream ?

Hash Functions

- Probabilistic data structures use hash functions to randomize items.
- They ignore collisions to keep the size constant
- But this is a reason why they can not give exact values.

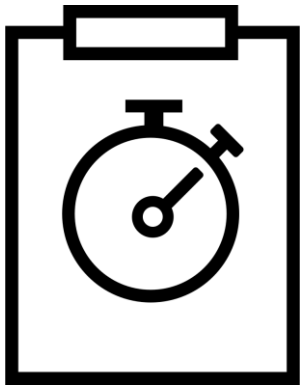

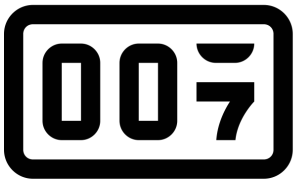

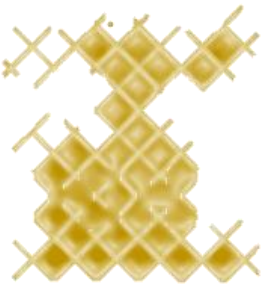

Hash Usage : Set Membership, Counting

SSN : Social Security Number

Is this guy on board ?

999999999999999

125675798988090

000000000000000

KEY

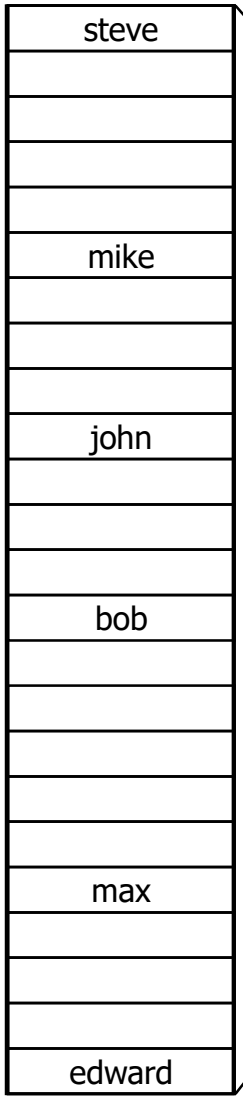

hash

N element Array

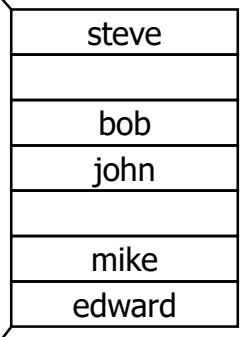

addresses

value

Hash Usage : Set Membership, Counting

SSN : Social Security Number

Is this guy on board ?

Probe O(1)

N element Array

addresses

steve	99
bob	
john	
mike	
edward	00

value

hash

steve
mike
john
bob
max
edward

999999999999999

125675798988090

KEY

000000000000000

Hash Usage : Set Membership, Counting

SSN : Social Security Number

How many of them are there ?

N element Array

Probe O(1)

addresses	value
	steve
	bob
	john
	mike
	edward

99

00

hash

	KEY
steve	999999999999999
mike	
john	125675798988090
bob	**KEY**
max	
edward	000000000000000

Hash Usage : Set Membership, Counting

SSN : Social Security Number

How many of them are there ?

N element Array

Probe O(1)

addresses		hash	value
steve		99	
bob			
john			
mike			
edward		00	

steve
mike
john
bob
max
edward

999999999999999

125675798988090

KEY

000000000000000

Hash Usage : Set Membership, Counting

SSN : Social Security Number

Is this guy on board ?

0	
1	
2	1
3	
4	
5	
6	1
7	
8	
9	
10	
11	1
12	
13	1
14	
15	

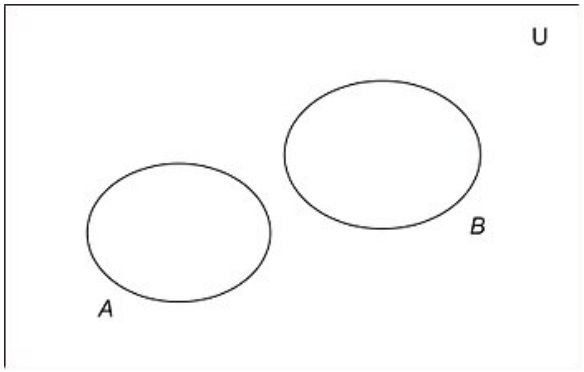

Hash Usage : Set Membership, Counting

SSN : Social Security Number

How many of them are there ?

0	3
1	
2	2
3	
4	5
5	
6	3
7	
8	
9	4
10	
11	1
12	
13	1
14	
15	2

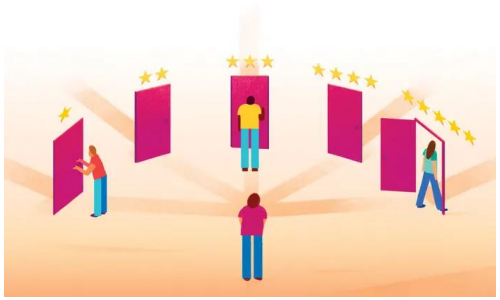

Hash Usage : Set Membership, Counting

SSN : Social Security Number

0	3
1	
2	2
3	
4	5
5	
6	3
7	
8	
9	4
10	
11	1
12	
13	1
14	
15	2

What about memory size ?

Exemple with alphabet letters

The set

A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z
1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26

Exemple with alphabet letters

The set

A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z
1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26

Bit Map : Probalistic Data Structure

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15

Exemple with alphabet letters

The set

A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z
1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26

Size = 27 * Short = 432 bytes

Probalistic Data Structure

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15

Size = 2 bytes

Hash set

Hash(A)

H(A)

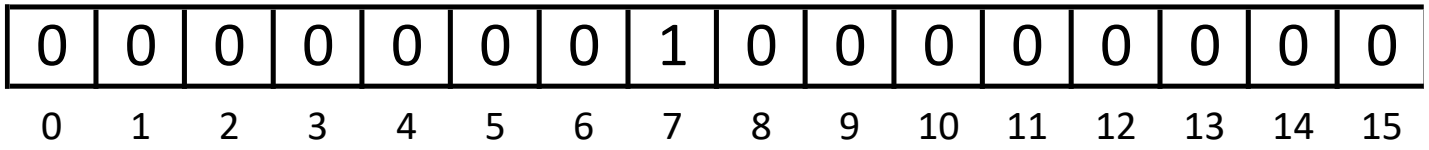

Hash set

Hash(B)

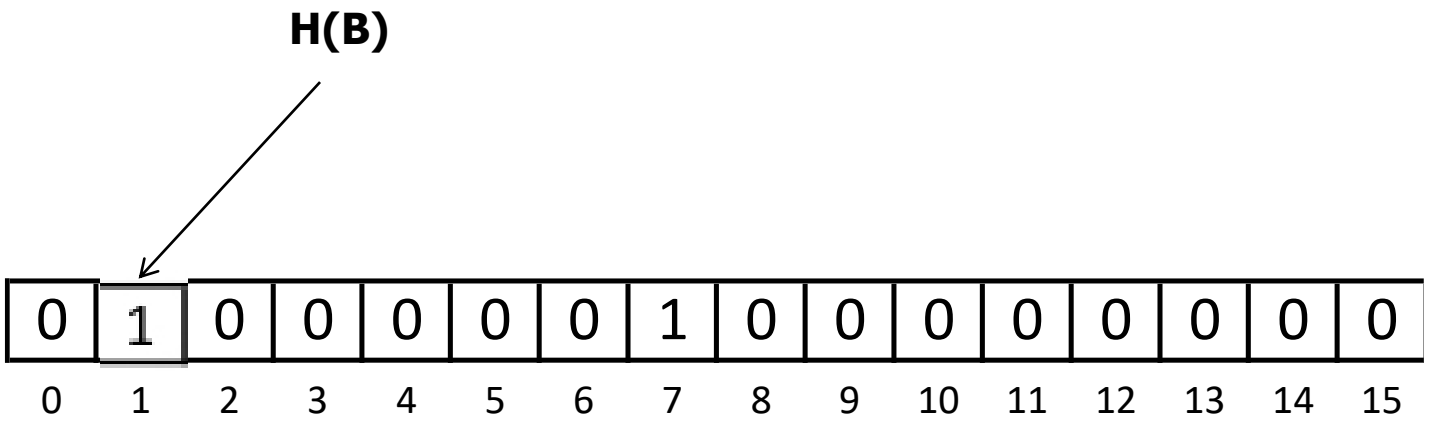

Hash set

Hash(Y)

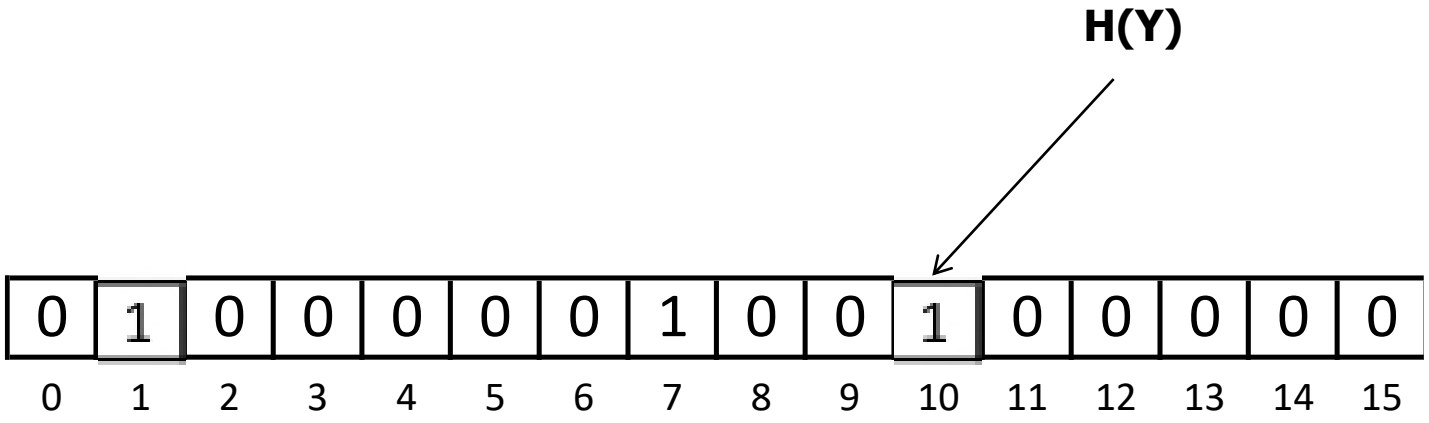

Hash set

Hash(E)= Hash(Y)

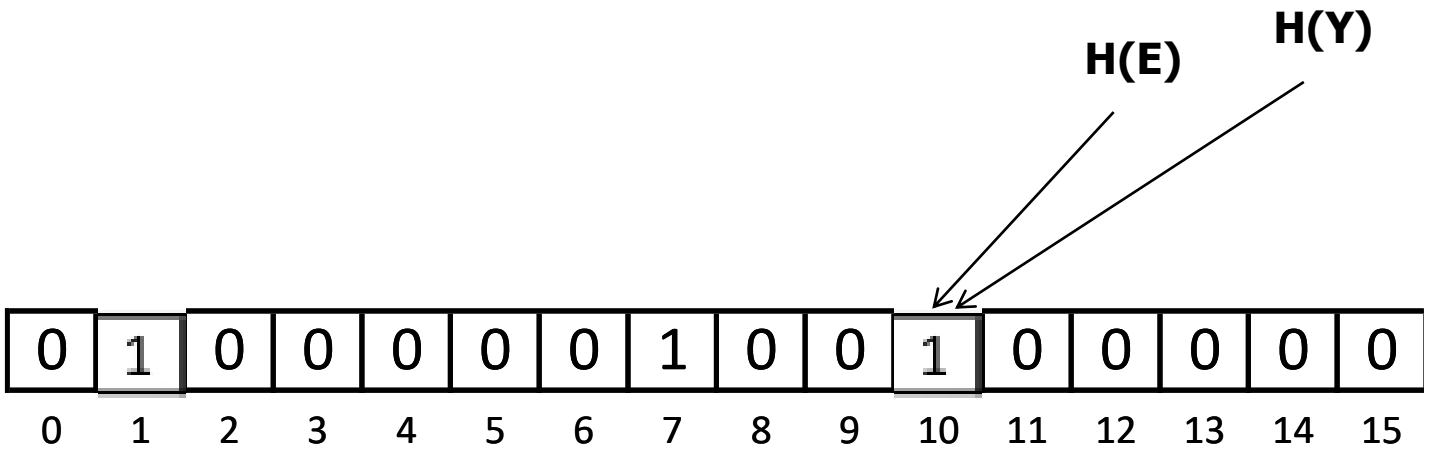

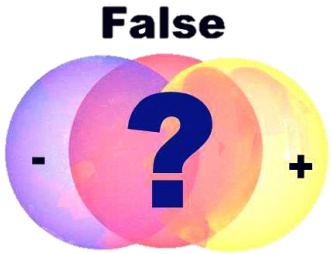

Hash set

Hash (K) = Hash (B)

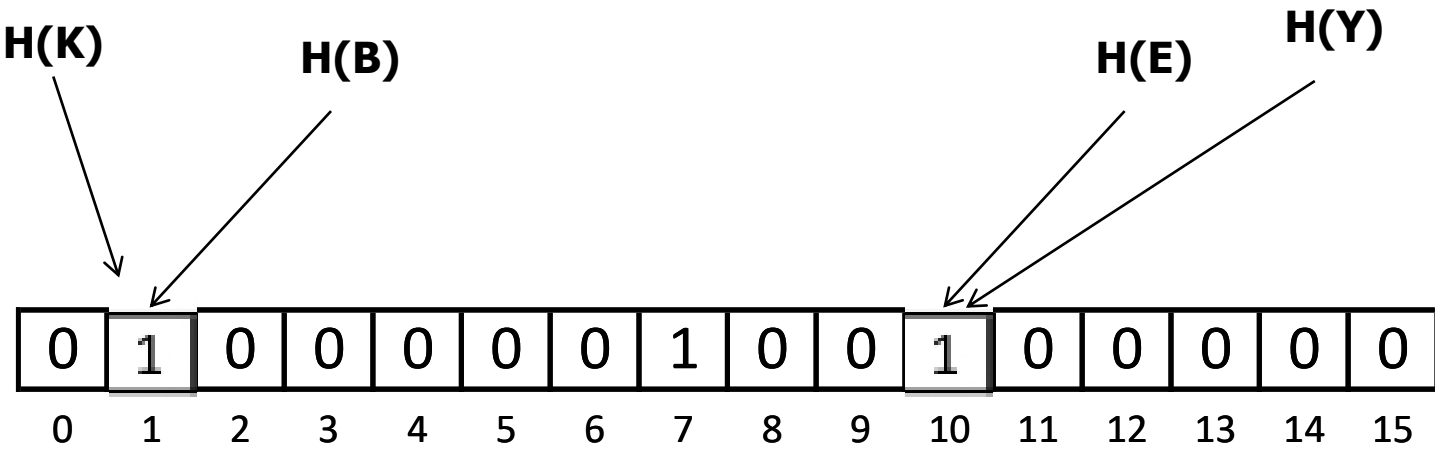

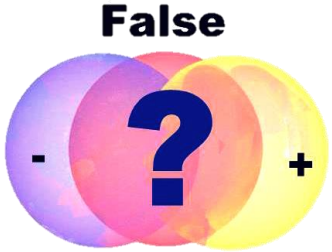

Fault Positive Prevention

- Two Hash functions : H1(), H2()
- Two Hash sets

0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15

0	1	0	0	0	0	0	1	0	0	0	0	0	0	0	0
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15

Fault Positive

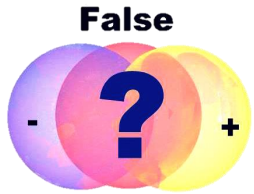

Collision

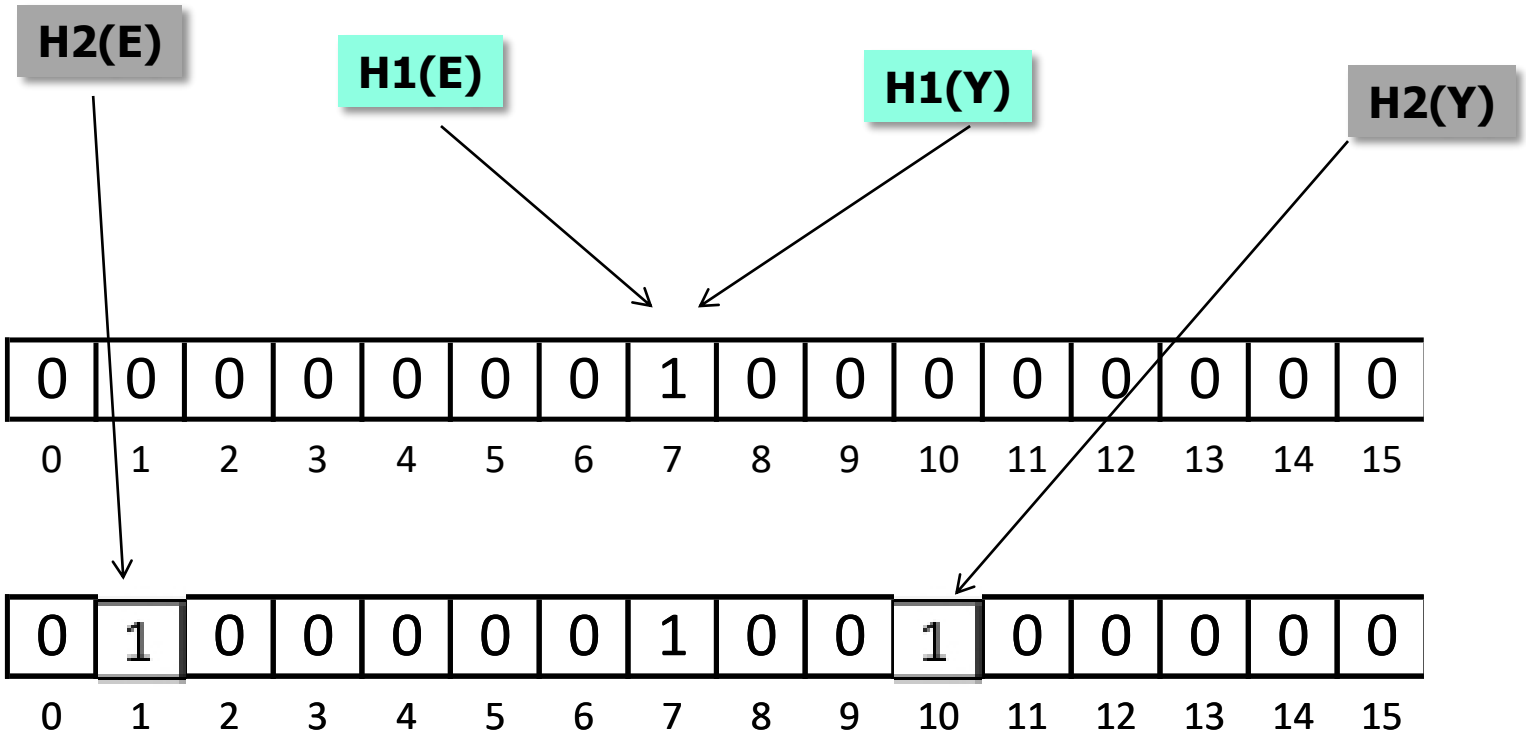

Fault Positive

Collision

If Y not In

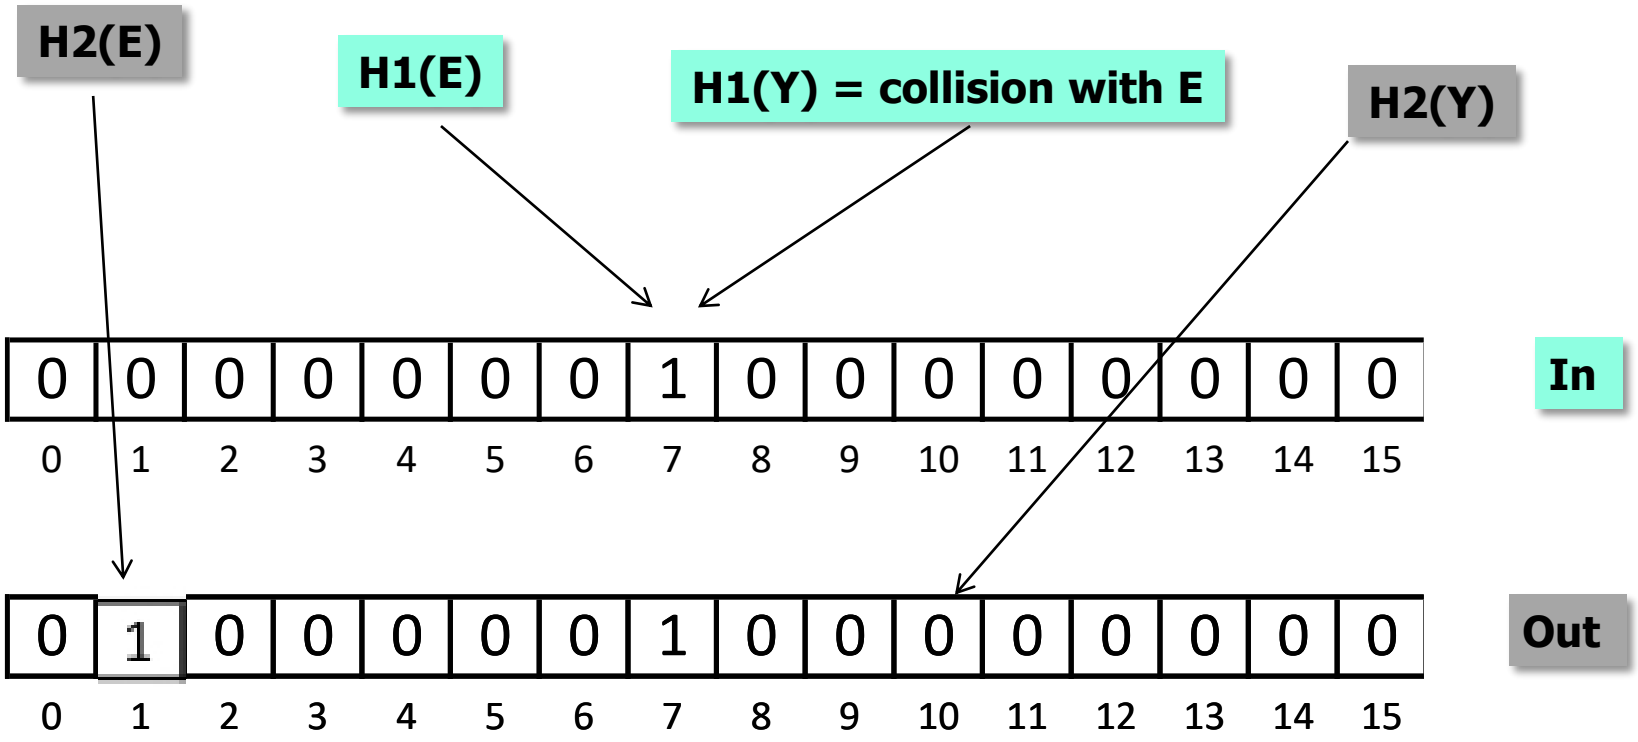

Y is definitiveley Out

Usual probabilistic data structures

- Bloom Filters
- Count-min Sketch
- Count-Mean-Min Sketch
- Hyperloglog


007